

# Gentle Introduction to Feedforward Neural Networks



**Heung-II Suk**

[hisuk@korea.ac.kr](mailto:hisuk@korea.ac.kr)

<http://www.ku-milab.org>



Department of Brain and Cognitive Engineering,  
Korea University

September 18-19, 2017



# Contents

- 1 Towards an Artificial Neural Network
- 2 Towards Feed-Forward Neural Networks
- 3 Network Training

# Towards an Artificial Neural Network

---

## Neuron: building block of the nervous system



[Image source: <https://askabiologist.asu.edu/neuron-anatomy>]

# Perceptron [Rosenblatt, 1958]

- Two-class model


- ▶ An input vector  $\mathbf{x}$  is first transformed using a fixed nonlinear transformation to give a feature vector  $\phi(\mathbf{x})$
- ▶ Then used to construct a **generalized linear model**

$$y(\mathbf{x}) = f(\mathbf{w}^\top \phi(\mathbf{x}))$$

$$\text{where } f(a) = \begin{cases} +1 & a \geq 0 \\ -1 & a < 0 \end{cases}$$

- Use a target coding scheme

- ▶  $t = +1$  for class  $C_1$ , and  $t = -1$  for  $C_2$
- ▶ Matching the choice of activation function



## [Parameters Learning]

### Error function minimization

- Error function: **number of misclassifications**
- This error function is a piecewise constant function of  $\mathbf{w}$  with discontinuities (c.f., regression)
- No closed-form solution (no derivatives exist for non-smooth functions)
- Take an iterative approach

# Perceptron Criterion

- Seeking  $\mathbf{w}$  such that

$$\left\{ \begin{array}{l} \mathbf{x}_n \in C_1 (t_n = +1) \text{ will have } \mathbf{w}^\top \phi(\mathbf{x}_n) \geq 0 \\ \mathbf{x}_n \in C_2 (t_n = -1) \text{ will have } \mathbf{w}^\top \phi(\mathbf{x}_n) < 0 \end{array} \right\} \Rightarrow \mathbf{w}^\top \phi(\mathbf{x}_n) t_n \geq 0$$

- ▶ Linearly bisecting the feature space

- For each misclassified sample, perceptron criterion tries to minimize

$$E_P(\mathbf{w}) = - \sum_{n \in M} \mathbf{w}^\top \phi(\mathbf{x}_n) t_n$$

$M$ : a set of all misclassified samples

# Perceptron Algorithm





- Error function:  $E_P(\mathbf{w}) = -\sum_{n \in M} \mathbf{w}^\top \phi(\mathbf{x}_n) t_n$
- Batch/Stochastic Gradient descent

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E_P(\mathbf{w}) = \mathbf{w}^{(\tau)} + \eta \sum_n \phi(\mathbf{x}_n) t_n$$

$\eta$ : learning rate,  $\tau$ : step index

- ▶ Since  $y(\mathbf{x}, \mathbf{w})$  is unchanged if we multiply  $\mathbf{w}$  by a constant, we can set  $\eta$  equal to 1 without loss of generality.
- Interpretation: cycle through the training samples in turn
  - ▶ If misclassified, for class  $C_1$  add  $\phi(\mathbf{x}_n)$  to  $\mathbf{w}$
  - ▶ If misclassified, for class  $C_2$  subtract  $\phi(\mathbf{x}_n)$  from  $\mathbf{w}$







Black arrow:  $w$  (points towards the decision region of the red class), green point: misclassified

# Changes in Non-linear Activation Function

1960's




1980's




# Hands on Programming: Perceptron Learning


# Rosenblatt's Perceptron with Vanilla Gradient




# Widrow-Hoff's Algorithm with Vanilla Gradient



# Widrow-Hoff's Algorithm with SDG



# Widrow-Hoff Perceptron with Minibatch-SDG




## **Towards Feed-Forward Neural Networks**

---




## (Recap.) Linear Basis Function Models - Perceptron

$$y(\mathbf{x}, \mathbf{w}) = f\left(\underbrace{\sum_{j=0}^D w_j x_j}_a\right)$$



- $f(\cdot)$ : continuous output
  - ▶ (regression) identity function; (classification) sigmoidal function

$$y(\mathbf{x}, \mathbf{w}) = f\left(\underbrace{\sum_{j=0}^M w_j \phi_j(\mathbf{x})}_a\right)$$




$\{\phi_j\}_{j=1}^M$ : basis functions

$$y(\mathbf{x}, \mathbf{w}) = f\left(\sum_{j=0}^M w_j \phi_j(\mathbf{x})\right)$$


For application to large-scale problems, it is necessary *to adapt the basis functions to the data*



$h(W\mathbf{x})$



Neural Networks




$K(\mathbf{x}, \mathbf{x}')$


Support Vector Machines

$$y(\mathbf{x}, \mathbf{w}) = f\left(\sum_{j=0}^M w_j \phi_j(\mathbf{x})\right)$$

- Basis functions  $\phi_j(\mathbf{x})$ : a parametric form
- These parameters to be adjusted along with the coefficients  $\{w_j\}$
- In NN, each basis function itself is a nonlinear function of a linear combination of the inputs
  - ▶ Coefficients in the linear combination are adaptive parameters




$\Rightarrow$



$h(\cdot)$ : nonlinear function

# Feed-Forward Network



$$\phi(\mathbf{x}) = h(W^{(1)}\mathbf{x}) = \mathbf{z}$$

$$W^{(1)} = [w_{ji}^{(1)}] \in \mathbb{R}^{M \times (D+1)}$$

$$\mathbf{y} = f(W^{(2)}\mathbf{z})$$

$$W^{(2)} = [w_{kj}^{(2)}] \in \mathbb{R}^{K \times (M+1)}$$

Two-layer neural network


- Pre-activations:  $\{a_j^{(1)}\}_{j=1}^M$

$$a_j^{(1)} = \sum_{i=1}^D \underbrace{w_{ji}^{(1)}}_{\text{weight}} x_i + \underbrace{w_{j0}^{(1)}}_{\text{bias}}$$

- Outputs of the basis functions ('hidden units')

$$z_j^{(1)} = h(a_j^{(1)})$$


- ▶  $h(\cdot)$ : differentiable, nonlinear activation function
- ▶ sigmoidal functions such as 'logistic sigmoid' or 'tanh'
- ▶ or others (discussed in deep learning)



Connection between input and hidden units

$$\text{sigmoid}(a) = \frac{1}{1 + \exp[-a]}$$

$$\tanh(a) = \frac{\exp[a] - \exp[-a]}{\exp[a] + \exp[-a]}$$




- Pre-activations:  $\{a_k^{(2)}\}_{k=1}^K$

$$a_k^{(2)} = \sum_{j=1}^M \underbrace{w_{kj}^{(2)}}_{\text{weight}} z_j^{(1)} + \underbrace{w_{k0}^{(2)}}_{\text{bias}}$$

- Outputs of the basis functions ('*output units*')  
 $y_k = f(a_k^{(2)})$

$$y_k = f(a_k^{(2)})$$



Connection between hidden and output units



- Choice of the output activation function  $f(\cdot)$  is determined by the nature of the data and the assumed distribution of target variables

- ▶ Regression: identity

$$y_k = a_k^{(2)}$$

- ▶ Binary classification: logistic sigmoid

$$y_k = \text{sigmoid} \left( a_k^{(2)} \right) = \frac{1}{1 + \exp \left( -a_k^{(2)} \right)}$$

- ▶ Multiclass classification: softmax

$$y_k = \text{softmax} \left( a_k^{(2)} \right) = \frac{\exp \left( a_k^{(2)} \right)}{\sum_{l=1}^K \exp \left( a_l^{(2)} \right)}$$

$$y_k(\mathbf{x}, \mathbf{w}) = f \left( \sum_{j=1}^M w_{kj}^{(2)} h \left( \sum_{i=1}^D w_{ji}^{(1)} x_i + w_{j0}^{(1)} \right) + w_{k0}^{(2)} \right)$$

equivalently

$$y_k(\mathbf{x}, \mathbf{w}) = f \left( \sum_{j=0}^M w_{kj}^{(2)} h \left( \underbrace{\sum_{i=0}^D w_{ji}^{(1)} x_i}_{\phi_j(\mathbf{x})} \right) \right)$$

where  $\mathbf{w} = \left[ \text{vec} \left( W^{(1)} \right); \text{vec} \left( W^{(2)} \right) \right]$

# Network Training

---

# Overview

- Error function
- Parameter learning: (stochastic) gradient-descent method
- Gradient evaluation: backpropagation [Rumelhart *et al.*, 1986]

# Error Functions

Given a set of input vectors  $\{\mathbf{x}_n\}_{n=1}^N$  and target vectors  $\{\mathbf{t}_n\}_{n=1}^N$

- Regression

- ▶ Identity activation function
- ▶ Sum-of-squares error

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \|\mathbf{y}(\mathbf{x}_n, \mathbf{w}) - \mathbf{t}_n\|^2$$

- Binary classification

- ▶ Logistic sigmoid activation function
- ▶ Cross-entropy error function

$$E(\mathbf{w}) = - \sum_{n=1}^N \{t_n \ln y_n + (1 - t_n) \ln (1 - t_n)\}$$


- Multiclass classification

- ▶ Softmax function
- ▶ Cross-entropy error function

$$E(\mathbf{w}) = - \sum_{n=1}^N \sum_{k=1}^K t_{kn} \ln y_{kn}$$

# Parameter Optimization

- Finding a weight vector  $\mathbf{w}$  that minimizes the error function  $E(\mathbf{w})$




- Geometrical picture of error function
  - ▶ a surface sitting over the weight space
- Small step from  $\mathbf{w}$  to  $\mathbf{w} + \Delta\mathbf{w}$  leads to change in error function

$$\Delta E \approx \Delta\mathbf{w}^\top \nabla E(\mathbf{w})$$

- ▶ At any point  $\mathbf{w}_C$ , the local gradient of the error surface is given by the vector

$$\nabla E = \left[ \frac{\partial E}{\partial w_1}, \frac{\partial E}{\partial w_2}, \dots, \frac{\partial E}{\partial w_d} \right]^\top$$



- $\nabla E$ : the direction of greatest rate of increase of the error function  $E(\mathbf{w})$
- To reduce the error, make a small step in the direction of  $-\nabla E(\mathbf{w})$
- Points at which the gradient vanishes: stationary points
  - minima, maxima, saddle points

- There is no analytical solution
- Resort to iterative numerical procedures
- Choose some initial value  $\mathbf{w}^{(0)}$  and then update it

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \Delta \mathbf{w}^{(\tau)}$$

- Different algorithms involve different choices for  $\Delta \mathbf{w}^{(\tau)}$ .
- Weight vector update  $\Delta \mathbf{w}^{(\tau)}$  is usually based on **gradient  $\nabla E(\mathbf{w})$  evaluated at the weight vector  $\mathbf{w}^{(\tau)}$**



## (Recap.) Gradient Descent Optimization

- Simplest approach to using gradient information
- Take a small step in the direction of the negative gradient

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E(\mathbf{w}^{(\tau)})$$

- ▶  $\eta$ : learning rate (e.g., 0.001)

$$E(\mathbf{w}) = \sum_{n=1}^N E_n(\mathbf{w})$$

## Different versions of gradient descent optimization

- **Batch**: the entire training set to be processes to evaluate  $\nabla E$   
(a.k.a., *gradient descent* or *steepest descent*)
- **On-line**: one sample at a time  
(a.k.a., *stochastic gradient descent* or *sequential gradient descent*)
  - ▶ Possibility of escaping from local minima
    - Stationary point w.r.t. the error function for the whole data set will generally not be a stationary point for each data point individually
- **Mini-batch stochastic gradient descent**: a small set of samples at a time
  - ▶ good tradeoff between batch and on-line
  - ▶ approximation of the gradient of the loss function

# Gradient Evaluation

To find an efficient technique for evaluating the gradient of an error function  $E(\mathbf{w})$  for a feed-forward neural network

- *Local message passing* scheme
- alternative information passing: forwards and backwards through the network

‘Error backpropagation’ or simply ‘Backprop’


## Simple Linear Model

$$y_k = \sum_i w_{ki} x_i$$
$$y_{nk} = y_k(\mathbf{x}_n, \mathbf{w})$$
$$E_n = \frac{1}{2} \sum_k (y_{nk} - t_{nk})^2$$

$$\frac{\partial E_n}{\partial w_{ki}} = \underbrace{(y_{nk} - t_{nk})}_{\text{error}} x_{ni}$$

- 'Local' computation

- ▶ an 'error signal'  $(y_{nk} - t_{nk})$  associated with the output end of the link  $w_{ki}$
- ▶ the variable  $x_{ni}$  associated with the input end of the link




## Forward Propagation in a General Feed-Forward Network

- 1 Each unit computes a weighted sum of its inputs

$$a_j^{(l)} = \sum_i w_{ji}^{(l)} z_i^{(l-1)}$$

- 2 Transformation by a nonlinear activation function  $h(\cdot)$  to give the activation  $z_j^{(l)}$  of unit  $j$

$$z_j^{(l)} = h(a_j^{(l)})$$



What is the error in this unit?

Evaluation of the derivative of  $E_n$  w.r.t.  $w_{ji}^{(l)}$

$$\begin{aligned}\frac{\partial E_n}{\partial w_{ji}^{(l)}} &= \underbrace{\frac{\partial E_n}{\partial a_j^{(l)}}}_{\equiv \delta_j^{(l)}} \underbrace{\frac{\partial a_j^{(l)}}{\partial w_{ji}^{(l)}}}_{z_i^{(l-1)}} && \text{(by chain rule)} \\ &= \delta_j^{(l)} z_i^{(l-1)}\end{aligned}$$

$\delta_j^{(l)}$ : gradient of the error function w.r.t. the pre-activation  $a_j^{(l)}$

- Note that this takes the same form as for the simple linear model
- Need only to calculate the value of  $\delta_j^{(l)}$  for each hidden and output unit in the network

- For hidden unit  $j$ ,


$$\delta_j^{(l)} = \frac{\partial E_n}{\partial a_j^{(l)}} = \sum_k \underbrace{\frac{\partial E_n}{\partial a_k^{(l+1)}}}_{\equiv \delta_k^{(l+1)}} \frac{\partial a_k^{(l+1)}}{\partial a_j^{(l)}}$$

- We are making use of the fact that variations in  $a_j^{(l)}$  give rise to variations in the error function only through variations in the variable  $a_k^{(l+1)}$

$$\frac{\partial a_k^{(l+1)}}{\partial a_j^{(l)}} = \frac{\partial \left( \sum_m w_{km}^{(l+1)} z_m^{(l)} \right)}{\partial a_j^{(l)}} = \frac{\partial \left( \sum_m w_{km}^{(l+1)} h \left( a_m^{(l)} \right) \right)}{\partial a_j^{(l)}} = h' \left( a_j^{(l)} \right) w_{kj}^{(l+1)}$$

$$\text{or } \frac{\partial a_k^{(l+1)}}{\partial a_j^{(l)}} = \frac{\partial a_k^{(l+1)}}{\partial z_j^{(l)}} \frac{\partial z_j^{(l)}}{\partial a_j^{(l)}} = w_{kj}^{(l+1)} h' \left( a_j^{(l)} \right)$$

$$\delta_j^{(l)} = h' \left( a_j^{(l)} \right) \sum_k w_{kj}^{(l+1)} \delta_k^{(l+1)}$$



Blue arrow: information flow during forward propagation  
 Red arrow: backward propagation of error information

The value of  $\delta$  for a particular hidden unit can be obtained by propagating the  $\delta$ 's backwards from units higher up in the network




- Forward propagation

$$a_j^{(l)} = \sum_i w_{ji}^{(l)} z_i^{(l-1)}$$
$$z_j^{(l)} = h\left(a_j^{(l)}\right)$$


- Backward propagation

$$\delta_j^{(l)} = h'\left(a_j^{(l)}\right) \sum_k w_{kj}^{(l+1)} \delta_k^{(l+1)}$$

## Forward propagation



## Backward propagation



$$\frac{\partial \mathbf{z}^{(1)}}{\partial \mathbf{a}^{(1)}} = h'(\mathbf{a}^{(1)})$$

$$\frac{\partial \mathbf{a}^{(2)}}{\partial \mathbf{z}^{(1)}} = W^{(2)}$$

$$\frac{\partial \mathbf{a}^{(1)}}{\partial W^{(1)}} = \mathbf{z}^{(0)} = \mathbf{x}$$

$$\frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{a}^{(2)}} = h'(\mathbf{a}^{(2)})$$


$$\frac{\partial \mathbf{a}^{(2)}}{\partial W^{(2)}} = \mathbf{z}^{(1)}$$

$$\frac{\partial E}{\partial \mathbf{z}^{(2)}} \equiv \delta^{(2)} = f'(\mathbf{a}^{(2)}) \odot (\mathbf{y} - \mathbf{t})$$


$$f'(\mathbf{a}) = \begin{cases} \frac{1}{1 - f(\mathbf{a})^2} & \text{logistic sigmoid} \\ \text{identity} \\ \text{tanh} \end{cases}$$




## Forward propagation




## Backward propagation




# Computational Graph




(a)



(b)

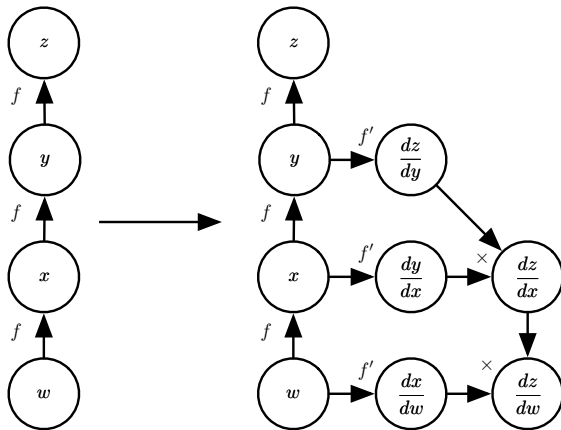




(c)

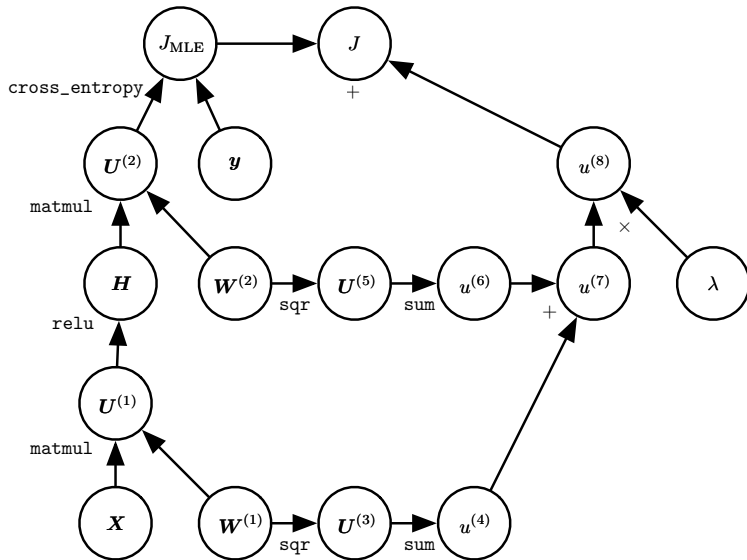


(d)

$$\begin{aligned}
 & \frac{\partial z}{\partial w} \\
 &= \frac{\partial z}{\partial y} \frac{\partial y}{\partial x} \frac{\partial x}{\partial w} \\
 &= f'(y) f'(x) f'(w) \\
 &= f'(f(f(w))) f'(f(w)) f'(w)
 \end{aligned}$$







# Hands on Programming: Neural Networks

---



**Thank you  
for your attention!!!**

**(Q & A)**

**hisuk (AT) korea.ac.kr**

<http://www.ku-milab.org>